1. Introduction

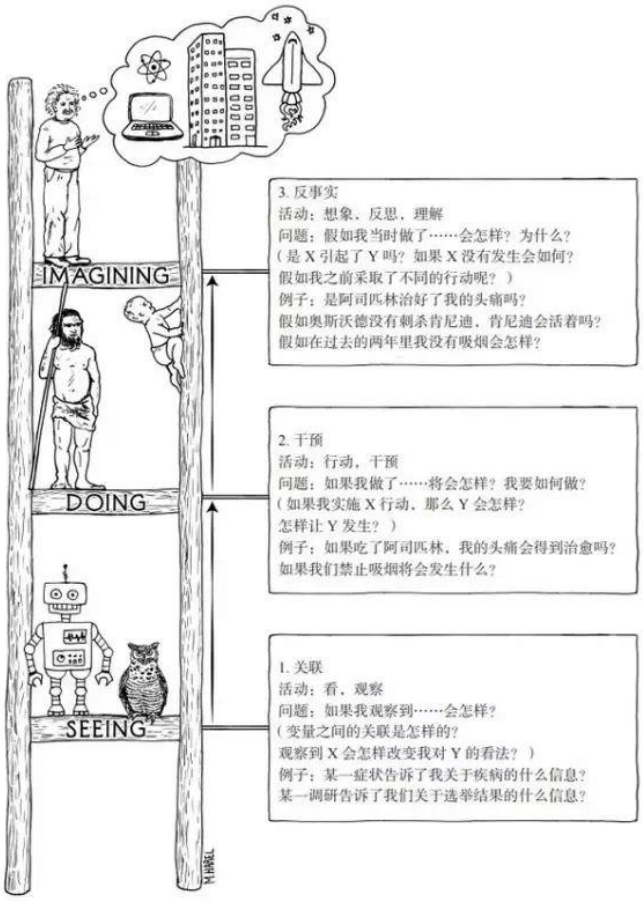

3. 反事实
活动：想象，反思，理解
问题：假如我当时做了……会怎样？为什么？
（是 X 引起了 Y 吗？如果 X 没有发生会如何？
假如我之前采取了不同的行动呢？）
例子：是阿司匹林治好了我的头痛吗？
假如奥斯沃德没有刺杀肯尼迪，肯尼迪会活着吗？
假如在过去的两年里我没有吸烟会怎样？

2. 干预
活动：行动，干预
问题：如果我做了……将会怎样？我要如何做？
（如果我实施 X 行动，那么 Y 会怎样？
怎样让 Y 发生？）
例子：如果吃了阿司匹林，我的头痛会得到治愈吗？
如果我们禁止吸烟将会发生什么？

1. 关联
活动：看，观察
问题：如果我观察到……会怎样？
（变量之间的关联是怎样的？
观察到 X 会怎样改变我对 Y 的看法？）
例子：某一症状告诉了我关于疾病的什么信息？
某一调研告诉了我们关于选举结果的什么信息？

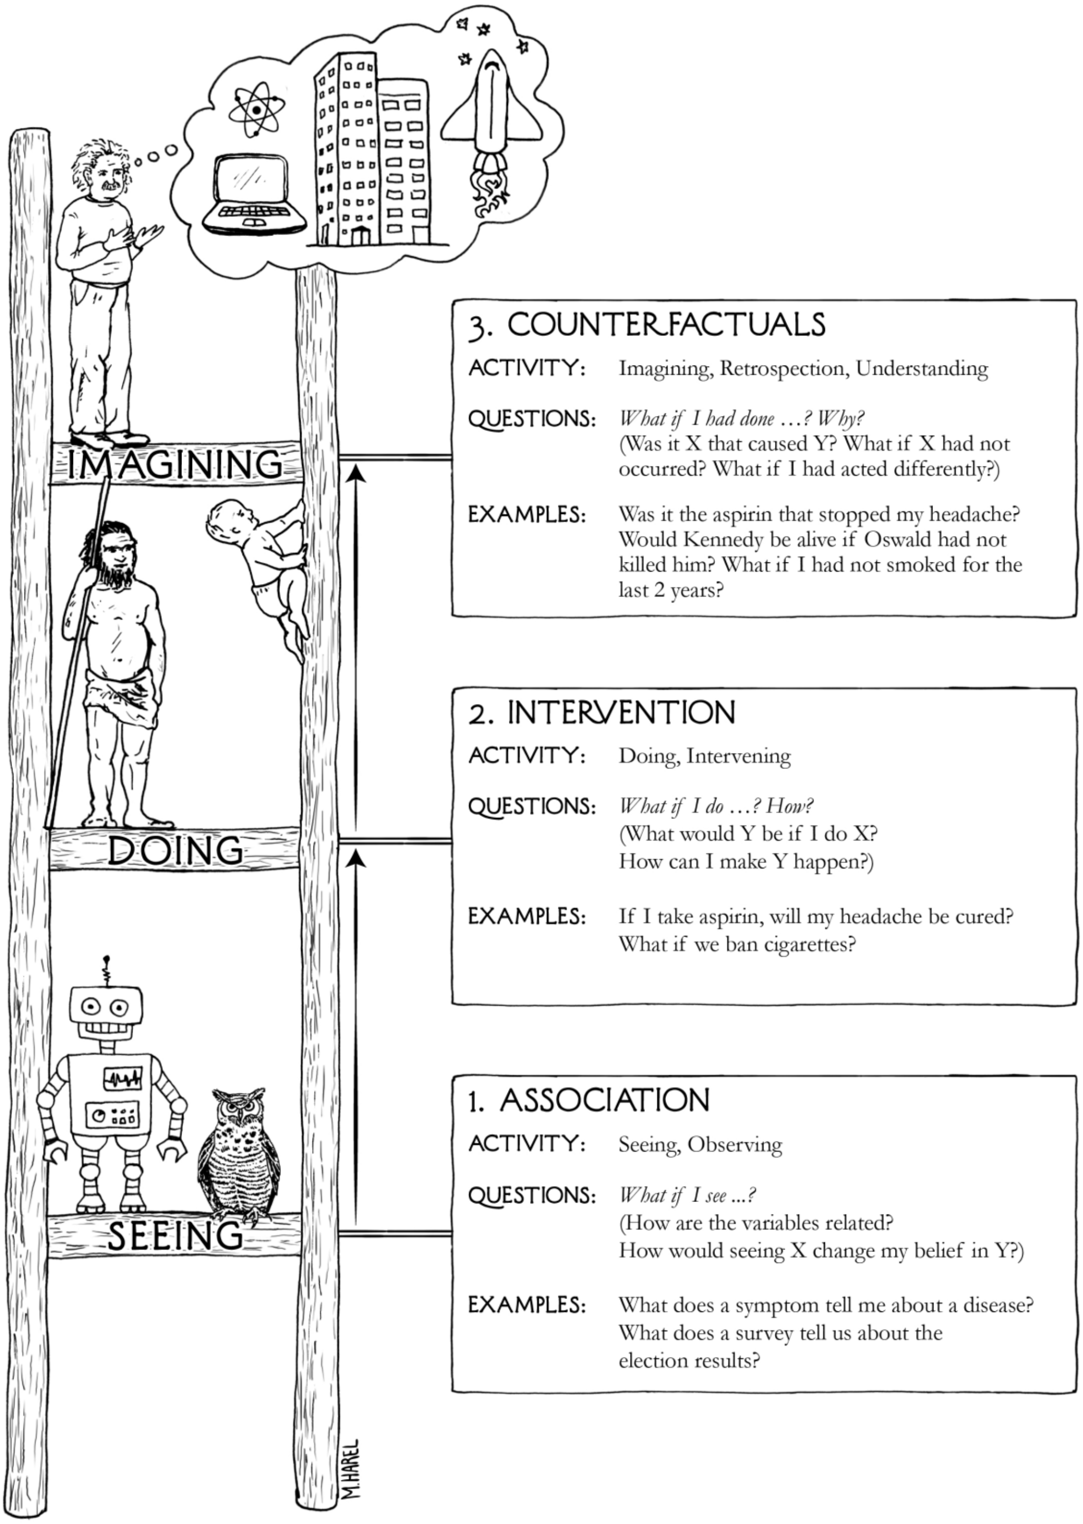

3. COUNTERFACTUALS

ACTIVITY: Imagining, Retrospection, Understanding

QUESTIONS: *What if I had done …? Why?*
(Was it X that caused Y? What if X had not occurred? What if I had acted differently?)

EXAMPLES: Was it the aspirin that stopped my headache? Would Kennedy be alive if Oswald had not killed him? What if I had not smoked for the last 2 years?

2. INTERVENTION

ACTIVITY: Doing, Intervening

QUESTIONS: *What if I do …? How?*
(What would Y be if I do X?
How can I make Y happen?)

EXAMPLES: If I take aspirin, will my headache be cured?
What if we ban cigarettes?

1. ASSOCIATION

ACTIVITY: Seeing, Observing

QUESTIONS: *What if I see ...?*
(How are the variables related?
How would seeing X change my belief in Y?)

EXAMPLES: What does a symptom tell me about a disease?
What does a survey tell us about the election results?

JUDEA PEARL

WINNER OF THE TURING AWARD

AND DANA MACKENZIE

THE BOOK OF WHY

α β

THE NEW SCIENCE OF CAUSE AND EFFECT

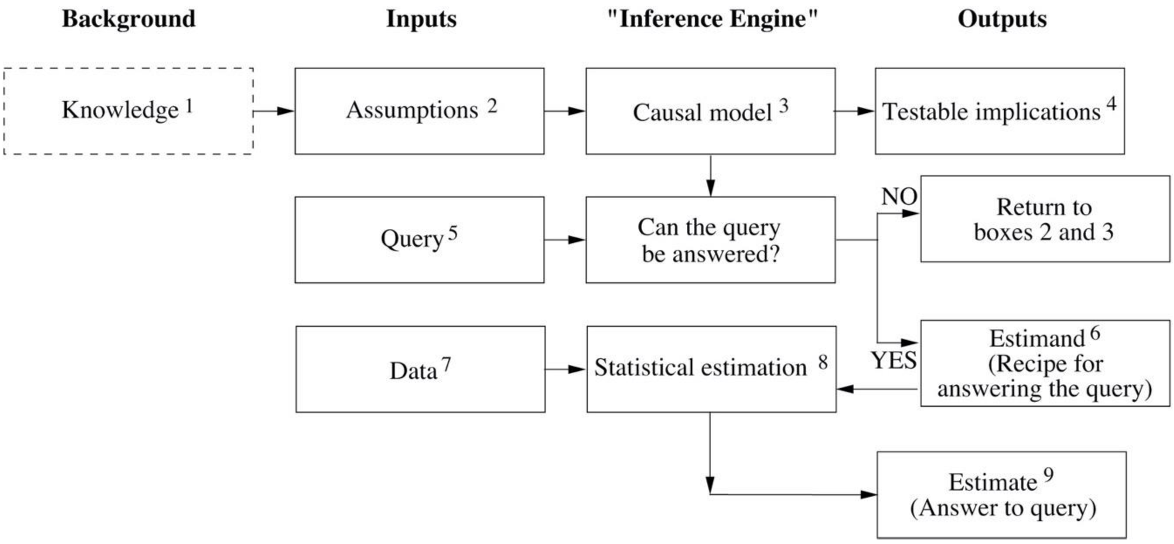
Background
Inputs
"Inference Engine"
Outputs
Knowledge 1
Assumptions 2
Causal model 3
Testable implications 4
Query 5
Can the query be answered?
NO
Return to boxes 2 and 3
YES
Estimand 6 (Recipe for answering the query)
Data 7
Statistical estimation 8
Estimate 9 (Answer to query)

2. Structural Causal Models

Structural causal models (SCMs) provide a comprehensive theory of causality.

2.1. Causal Graph

A causal graph $G = (V, E)$ is a directed graph that describes the causal effects between variables, where V is the node set and E the edge set. In a causal graph, each node represents a random variable including the treatment, the outcome, other observed and unobserved variables. A directed edge $x \rightarrow y$ denotes a causal effect of x on y.

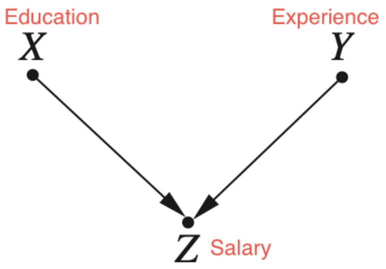
Education
X
Experience
Y
Z
Salary

2.2. Product Decomposition

For any model whose graph is acyclic, the joint distribution of the variables in the model is given by the product of the conditional distributions $P(child|parents)$ over all the “families” in the graph. Formally, we write this rule as

$$P(x_1, x_2, \dots, x_n) = \prod_i P(x_i | pa_i)$$

where pa_i stands for the values of the parents of variable X_i.

For example, in a simple chain graph $X \rightarrow Y \rightarrow Z$, we can write directly:

$$P(X = x, Y = y, Z = z) = P(X = x)P(Y = y | X = x)P(Z = z | Y = y)$$

This knowledge allows us to save an enormous amount of space when laying out a joint distribution. To estimate the joint distribution from a data set generated by the above model, we need not count the frequency of every triple (x, y, z); we can instead count the frequencies of each x, $(y|x)$, and $(z|y)$ and multiply.

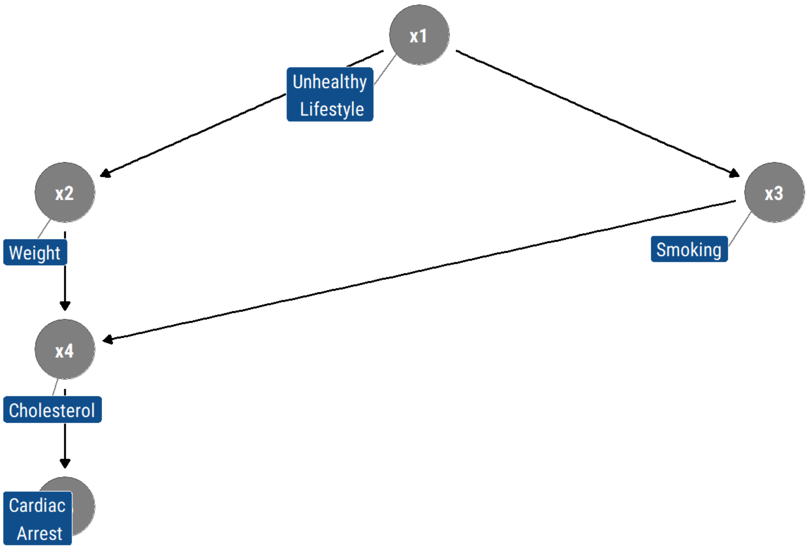

$$P(x_1, x_2, x_3, x_4, x_5) = P(x_1)P(x_2|x_1)P(x_3|x_1)P(x_4|x_2, x_3)P(x_5|x_4)$$

2.3. $\boldsymbol{d}$-Separation

Given a SCM, the conditional independence embedded in its causal graph provides sufficient information to confirm whether it satisfies the criteria such that we can apply certain causal inference methods. To understand the **conditional independence**, we need the concept of **dependency-separation** (d-separation) based on the definition of **blocked** path.

Two events A and B are conditionally independent given a third event C if $P(A|B, C) = P(A|C)$ and $P(B|A, C) = P(B|C)$.

The probabilistic implications of the d-separation criterion:

- If X and Y are d-separated by Z in causal graph G, then X is independent of Y

conditional on Z in every distribution compatible with G

- Conversely, if X and Y are not d-separated by Z in causal graph G, then X and Y are dependent conditional on Z

There are three typical DAGs for conditional independence:

- **Chain**: x causally affects y through its influence on z
- **Fork**: z is the common cause of both x and y, that is, x is associated with y but there is no causation between them
- **Collider**: both x and y cause z but there is no causal effect or association between x and y

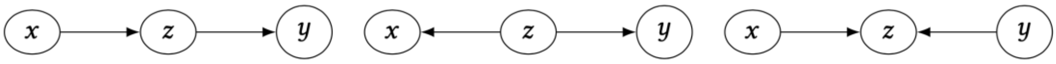

(a) Chain (b) Fork (c) Collider

In the chain and fork, the path between x and y is blocked, that is, the variables become independent, if we condition on z. Contrarily, in a collider, conditioning on z introduces an association between x and y. Generally, we say conditioning on a set of nodes $\mathcal{Z}$ blocks a path p iff there exists at least one node $z \in \mathcal{Z}$ in the path p.

Definition 1 ($\boldsymbol{d}$-Separation)

A path p is blocked (or d-separated) by a set of nodes $\mathcal{Z}$ if and only if

1. p contains a chain of nodes $A \to B \to C$ or a fork $A \leftarrow B \to C$ such that the middle node B is in $\mathcal{Z}$ (i.e., B is conditioned on), or
2. p contains a collider $A \leftarrow B \to C$ such that the collision node B is not in $\mathcal{Z}$, and no descendant of B is in $\mathcal{Z}$.

Chain

Rule 1 (Conditional Independence in Chains)

Two variables, X and Y, are conditionally independent given Z, if there is only one unidirectional path between X and Y, and Z is any set of variables that intercepts that path.

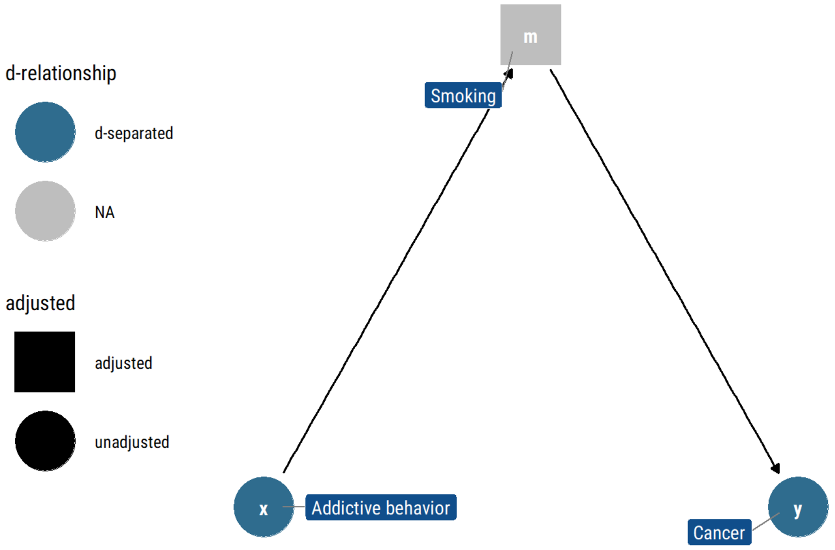
d-relationship
d-separated
NA
adjusted
adjusted
unadjusted
m
Smoking
x
Addictive behavior
y
Cancer

Addictive behavior and cancer are d-separated by smoking

Fork

Rule 2 (Conditional Independence in Forks)

If a variable X is a common cause of variables Y and Z, and there is only one path between Y and Z, then Y and Z are independent conditional on X.

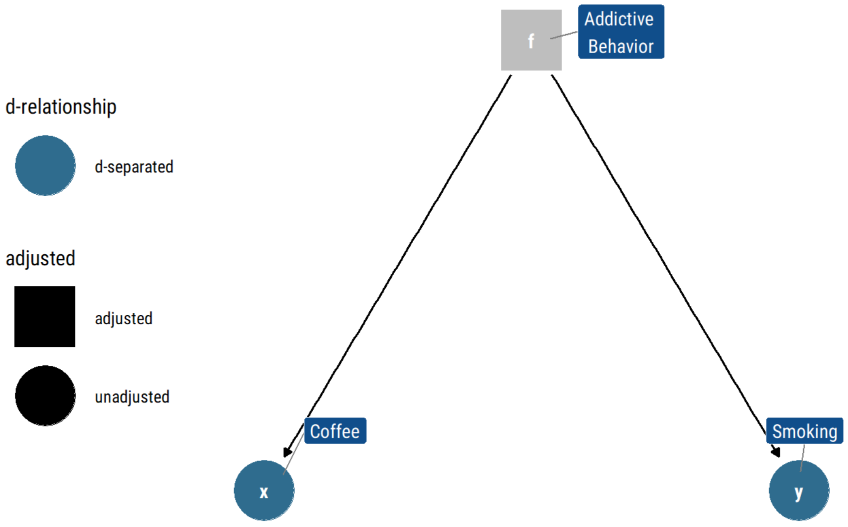
d-relationship
d-separated
adjusted
adjusted
unadjusted
f
Addictive Behavior
Coffee
x
Smoking
y

Drinking coffee and Smoking are d-separated by addictive behavior

Collider

Rule 3 (Conditional Independence in Colliders)

If a variable Z is the collision node between two variables X and Y, and there is only one path between X and Y, then X and Y are unconditionally independent but are dependent conditional on Z and any descendants of Z.

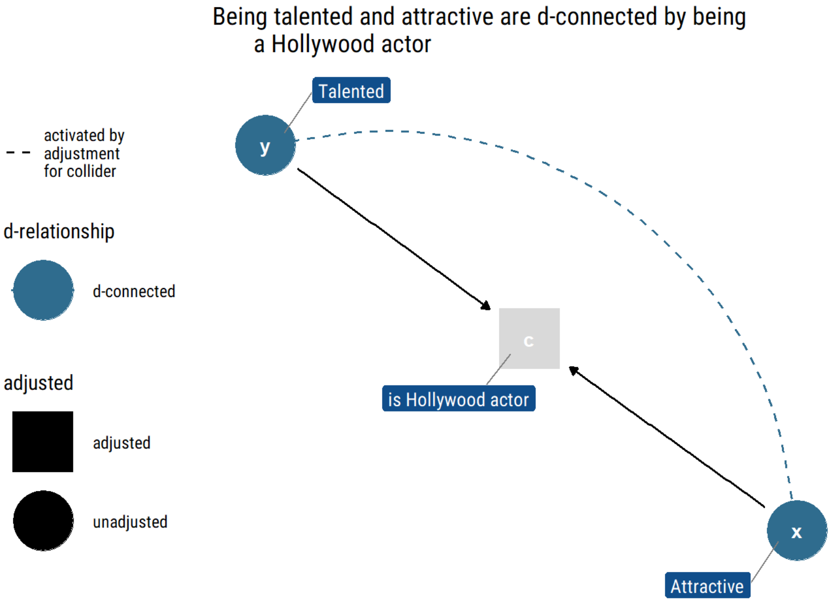

Being talented and attractive are d-connected by being a Hollywood actor

Talented
y
activated by adjustment for collider
d-relationship
d-connected
adjusted
adjusted
unadjusted
c
is Hollywood actor
x
Attractive

Example

Armed with the tool of d-separation, we can now look at some more complex graphical models and determine which variables in them are independent and dependent, both marginally and conditional on other variables.

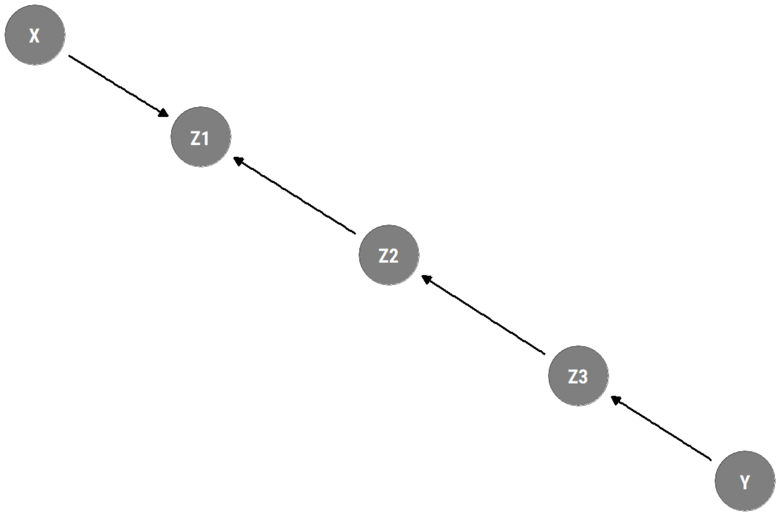

On which variable should we condition to make X and Y conditionally dependent?

Check whether the path is opened. If not, open it using d-separation criterion

Can we distinguish models from data alone?

The preceding sections demonstrate that causal models have **testable implications** in the data sets they generate. For instance, if we have a graph G that we believe might have generated a data set S, d-separation will tell us which variables in G must be independent conditional on which other variables. Conditional independence is something we can test for using a data set. Suppose we list the d-separation conditions in G, and note that variables A and B must be independent conditional on C. Then, suppose we estimate the probabilities based on S, and discover that the data suggests that A and B are not independent conditional on C. We can then reject G as a possible causal model for S.

But there are other models that imply the same conditional independencies. Two graphical models G_1 and G_2 are observationally equivalent when they imply the same conditional independencies. The set of all the models with indistinguishable implications is called an **equivalence class**.

If two models are observationally equivalent, we **cannot use data alone to distinguish from them**. We must use our structural knowledge about the problem at hand to decide which model is the right one.

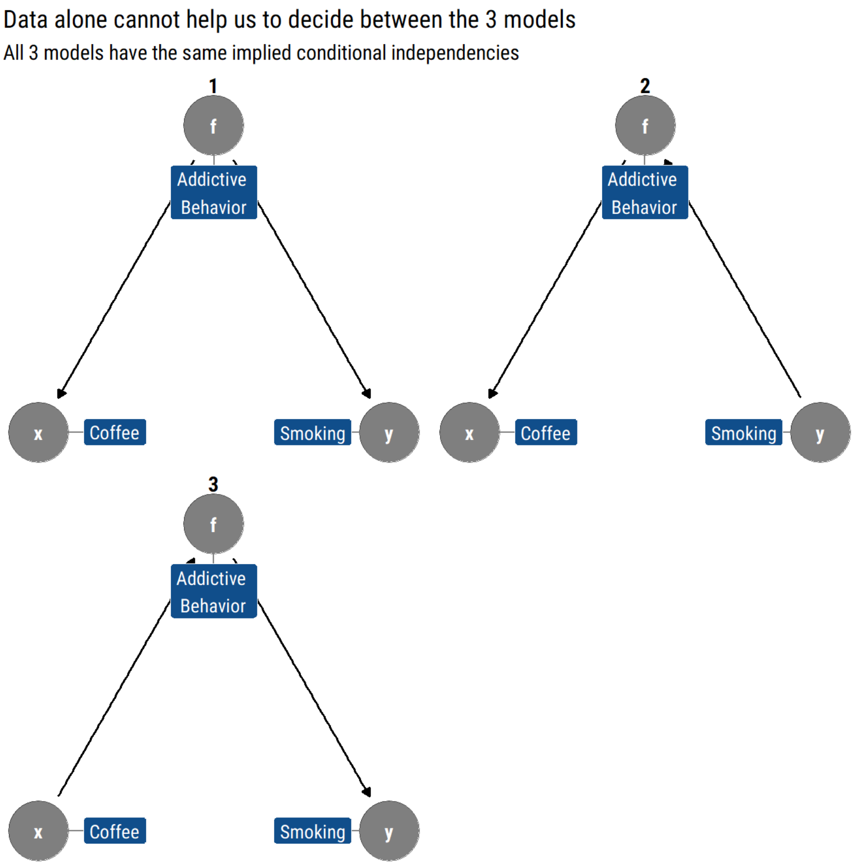
Data alone cannot help us to decide between the 3 models

All 3 models have the same implied conditional independencies

As Pearl says, **data are fundamentally dumb**: if we rely only in data to inform our models, we are extremely limited on what we can learn from them. We cannot predict the consequences of intervening (i.e., causal effects) in one of the variables with only observational data. Therefore, we must extend our theory beyond conditional probabilities gain causal understanding with only observational data.

3. The Effects of Interventions

Causal effects cannot be estimated from the data itself without a causal story, but this dost not mean that we cannot gain any causal understanding, far from which it just means that we must have a causal model. With the combination of causal models and observational data, causal effects can be estimated by leveraging the invariant information from the pre-intervention distribution.

3.1. Interventions

The **do-calculus** of Pearl was proposed to define **intervention** in SCMs. Specifically, the do-calculus introduces a new operator $do(t)$, which denotes the intervention of setting the value of the variable t to t'. The notation of $do(t)$ leads to a formal expression of the interventional distributions:

Definition 2 (Interventional Distribution)

The interventional distribution $P(y|do(x'))$ denotes the distribution of the variable y when we rerun the modified data-generation process where the value of variable x is set to x'.

Consider a study of how Yelp ratings influence potential restaurant customers. Below are SCMs without and under intervention $do(t')$ for the study, where x, t and y denote restaurant category, Yelp rating and customer flow. For the causal graph, the interventional distribution $P(y|do(t'))$ refers to the distribution of variable y as if the rating t is set to t' by intervention, **where all the arrows into t are removed**.

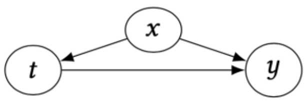

(a) An SCM without intervention.

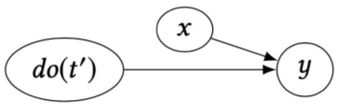

(b) An SCM under the intervention $do(t)$.

The difference between **intervening** on a variable and **conditioning** on that variable should, hopefully, be obvious. When we intervene on a variable in a model, we fix its value. We change the system, and the values of other variables often change as a result. When we condition on a variable, we change nothing; we merely narrow our focus to the subset of cases in which the variable takes the value we are interested in. What changes, then, is our perception about the world, not the world itself.

In notation, we distinguish between cases where a variable X takes a value x naturally and cases where we fix $X = x$ by denoting the latter $do(X = x)$. So $P(Y = y|X = x)$ is the probability that $Y = y$ conditional on finding $X = x$, while $P(Y = y|do(X = x))$ is the probability that $Y = y$ when we intervene to make $X = x$. In the distributional terminology, $P(Y = y|X = x)$ reflects the population distribution of Y among individuals whose X value is x. On the other hand, $P(Y = y|do(X = x))$ represents the population distribution of Y if **everyone in the population had their X value fixed at x**. We similarly write $P(Y = y|do(X = x), Z = z)$ to denote the conditional probability of $Y = y$, given $Z = z$, in the distribution created by the intervention $do(X = x)$.

3.2. The Adjustment Formula

In the classical example of Simpson's paradox, where X stands for drug usage, Y stands for recovery, and Z stands for gender. To find out how effective the drug is in the population, we imagine a hypothetical intervention by which we administer the drug uniformly to the entire population and compare the recovery rate to what would obtain under the complementary intervention, where we prevent everyone from using the drug. Denoting the first intervention by $do(X = 1)$ and the second by $do(X = 0)$, our task is to estimate the difference

$$P\big(Y = 1\big|do(X = 1)\big) - P\big(Y = 1\big|do(X = 0)\big)$$

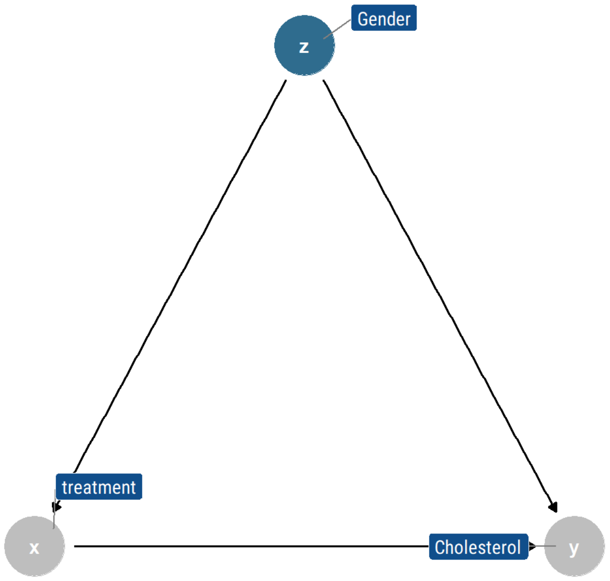

We know from first principles that causal effects cannot be estimated from the data set itself without a causal story. That was the lesson of Simpson’s paradox: The data itself was not sufficient even for determining whether the effect of the drug was positive or negative. But with the aid of the graph G above, we can compute the magnitude of the causal effect from the data. To do so, we simulate the intervention in the form of a graph surgery (showed in the graph G_m below), then the causal effect $P\left(Y = y \middle| do(X = x)\right)$ is equal to the conditional probability $P_m(Y = y|X = x)$ that prevails in the **manipulated** model, where P is the pre-intervention distribution probability and P_m is the post-intervention distribution probability.

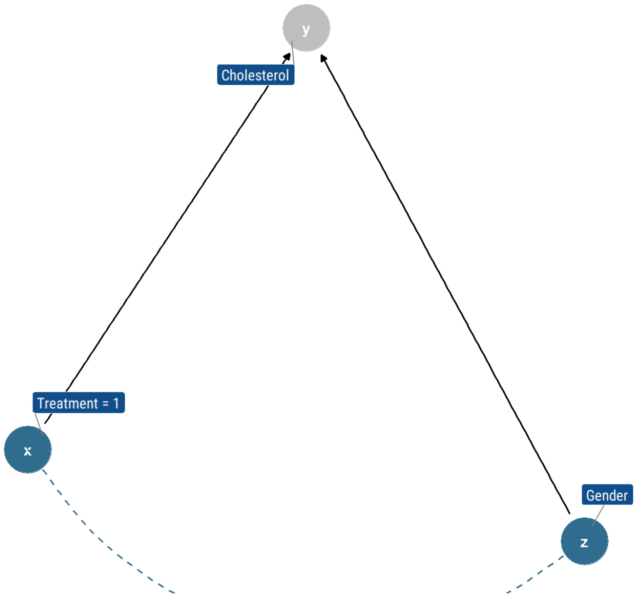

The key to computing the causal effect lies in the observation that P_m, the manipulated probability, shares two essential properties with P (the original probability function that prevails in the pre-intervention model):

- First, the marginal probability $P(Z = z)$ is invariant under the intervention, because the process determining Z is not affected by removing the arrow from Z to X
- Second, the conditional probability $P(Y = y|Z = z, X = x)$ is invariant, because the process by which Y responds to X and Z remains the same, regardless of whether X changes spontaneously or by deliberate manipulation

Therefore, two of the equations of invariance are $P_m(Y = y|Z = z, X = x) = P(Y = y|Z = z, X = x)$ and $P_m(Z = z) = P(Z = z)$.

We can also use the fact that Z and X are d-separated in the modified model and are, therefore, independent under the intervention distribution, which tells us that $P_m(Z = z|X = x) = P_m(Z = z) = P(Z = z)$. Putting these considerations together, we have

$$\begin{aligned} P\big(Y = y \mid do(X = x)\big) &= P_m(Y = y \mid X = x) \quad \text{(by definition)} \\ &= \sum_z P_m(Y = y \mid X = x, Z = z)\, P_m(Z = z \mid X = x) \\ &= \sum_z P_m(Y = y \mid X = x, Z = z)\, P_m(Z = z) \end{aligned}$$

Finally, using the invariance relations, we obtain a formula for the causal effect, in terms of pre-intervention probabilities:

$$P\big(Y = y \mid do(X = x)\big) = \sum_z P_m(Y = y \mid X = x, Z = z)\, P_m(Z = z)$$

which is called *adjustment formula*, and it computes the association between X and Y for each value z of Z, then averages over those values. This procedure is referred as “adjusting for Z” or “controlling for Z.”