

Code: <https://github.com/microsoft/dowhy>

Paper: <https://arxiv.org/abs/2011.04216>



DoWhy: An End-to-End Library for Causal Inference

Amit Sharma (@amt_shrma), Emre Kiciman (@emrek)

Microsoft Research

Causal ML group @ Microsoft:

<https://www.microsoft.com/en-us/research/group/causal-inference/>

From prediction to decision-making



Decision-making: Acting/intervening based on analysis

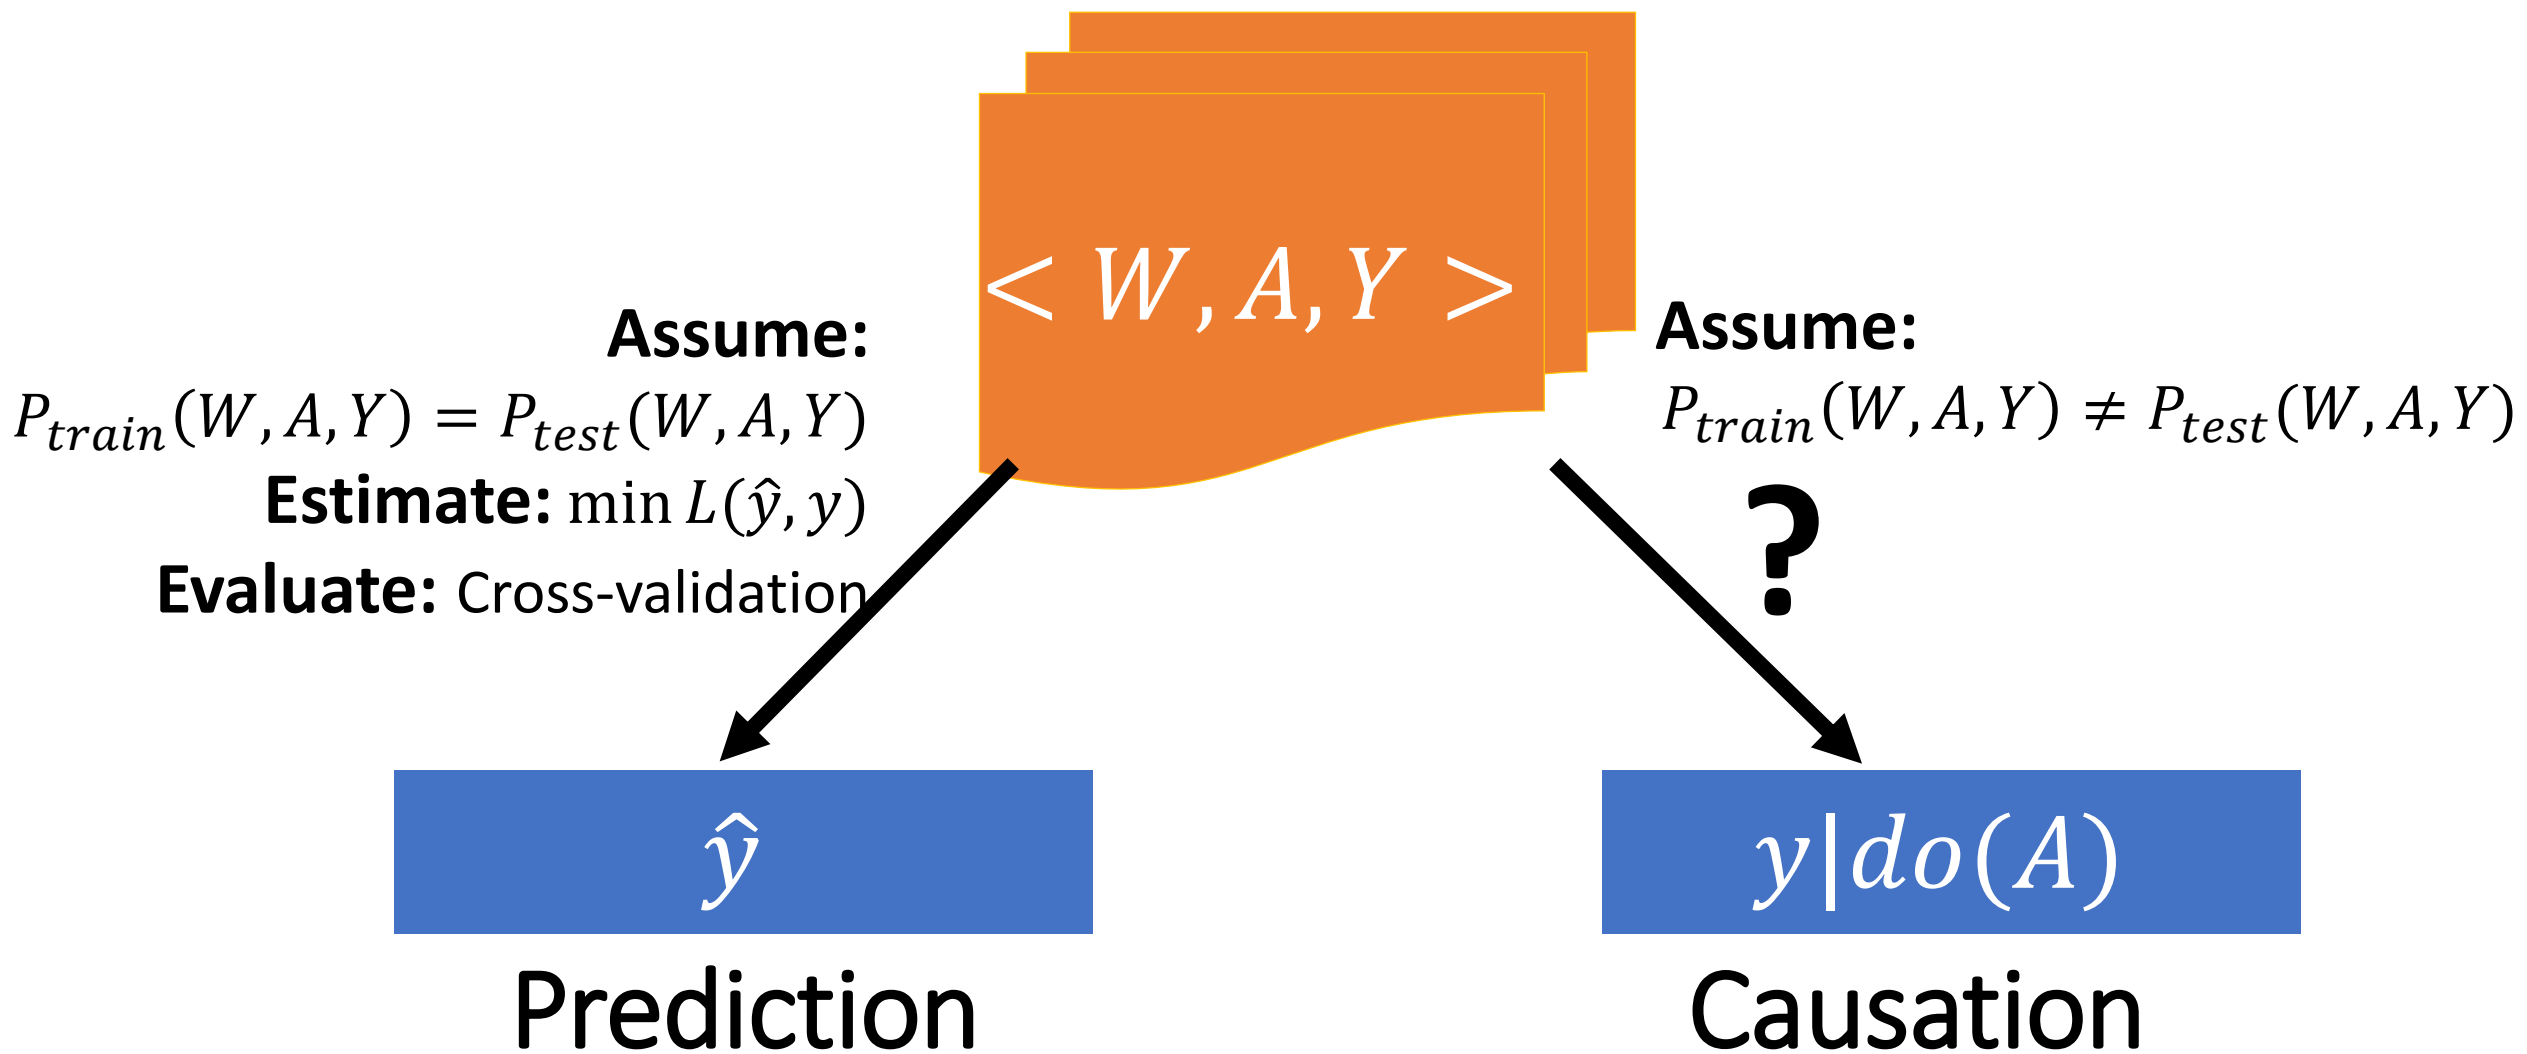
- Interventions break correlations used by conventional ML
- The feature with the highest importance score in a prediction model,
 - Need not be the best feature to act on
 - May not even affect the outcome at all!

For decision-making, need to find the features that **cause the outcome** & *estimate how the outcome would change if the features are changed.*

Many data science Qs are causal Qs

- **A/B experiments:** If I change the algorithm, will it lead to a higher success rate?
- **Policy decisions:** If we adopt this treatment/policy, will it lead to a healthier patient/more revenue/etc.?
- **Policy evaluation:** knowing what I know now, did my policy help or hurt?
- **Credit attribution:** are people buying because of the recommendation algorithm? Would they have bought anyway?

In fact, causal questions form the basis of almost all scientific inquiry



Two Fundamental Challenges for Causal Inference

Multiple causal mechanisms and estimates can fit the same data distribution.



1. Assumptions

Estimation about different data distributions than the training distribution (no easy “cross-validation”).



2. Evaluation



Real World: **do(A=1)**



Counterfactual World: **do(A=0)**

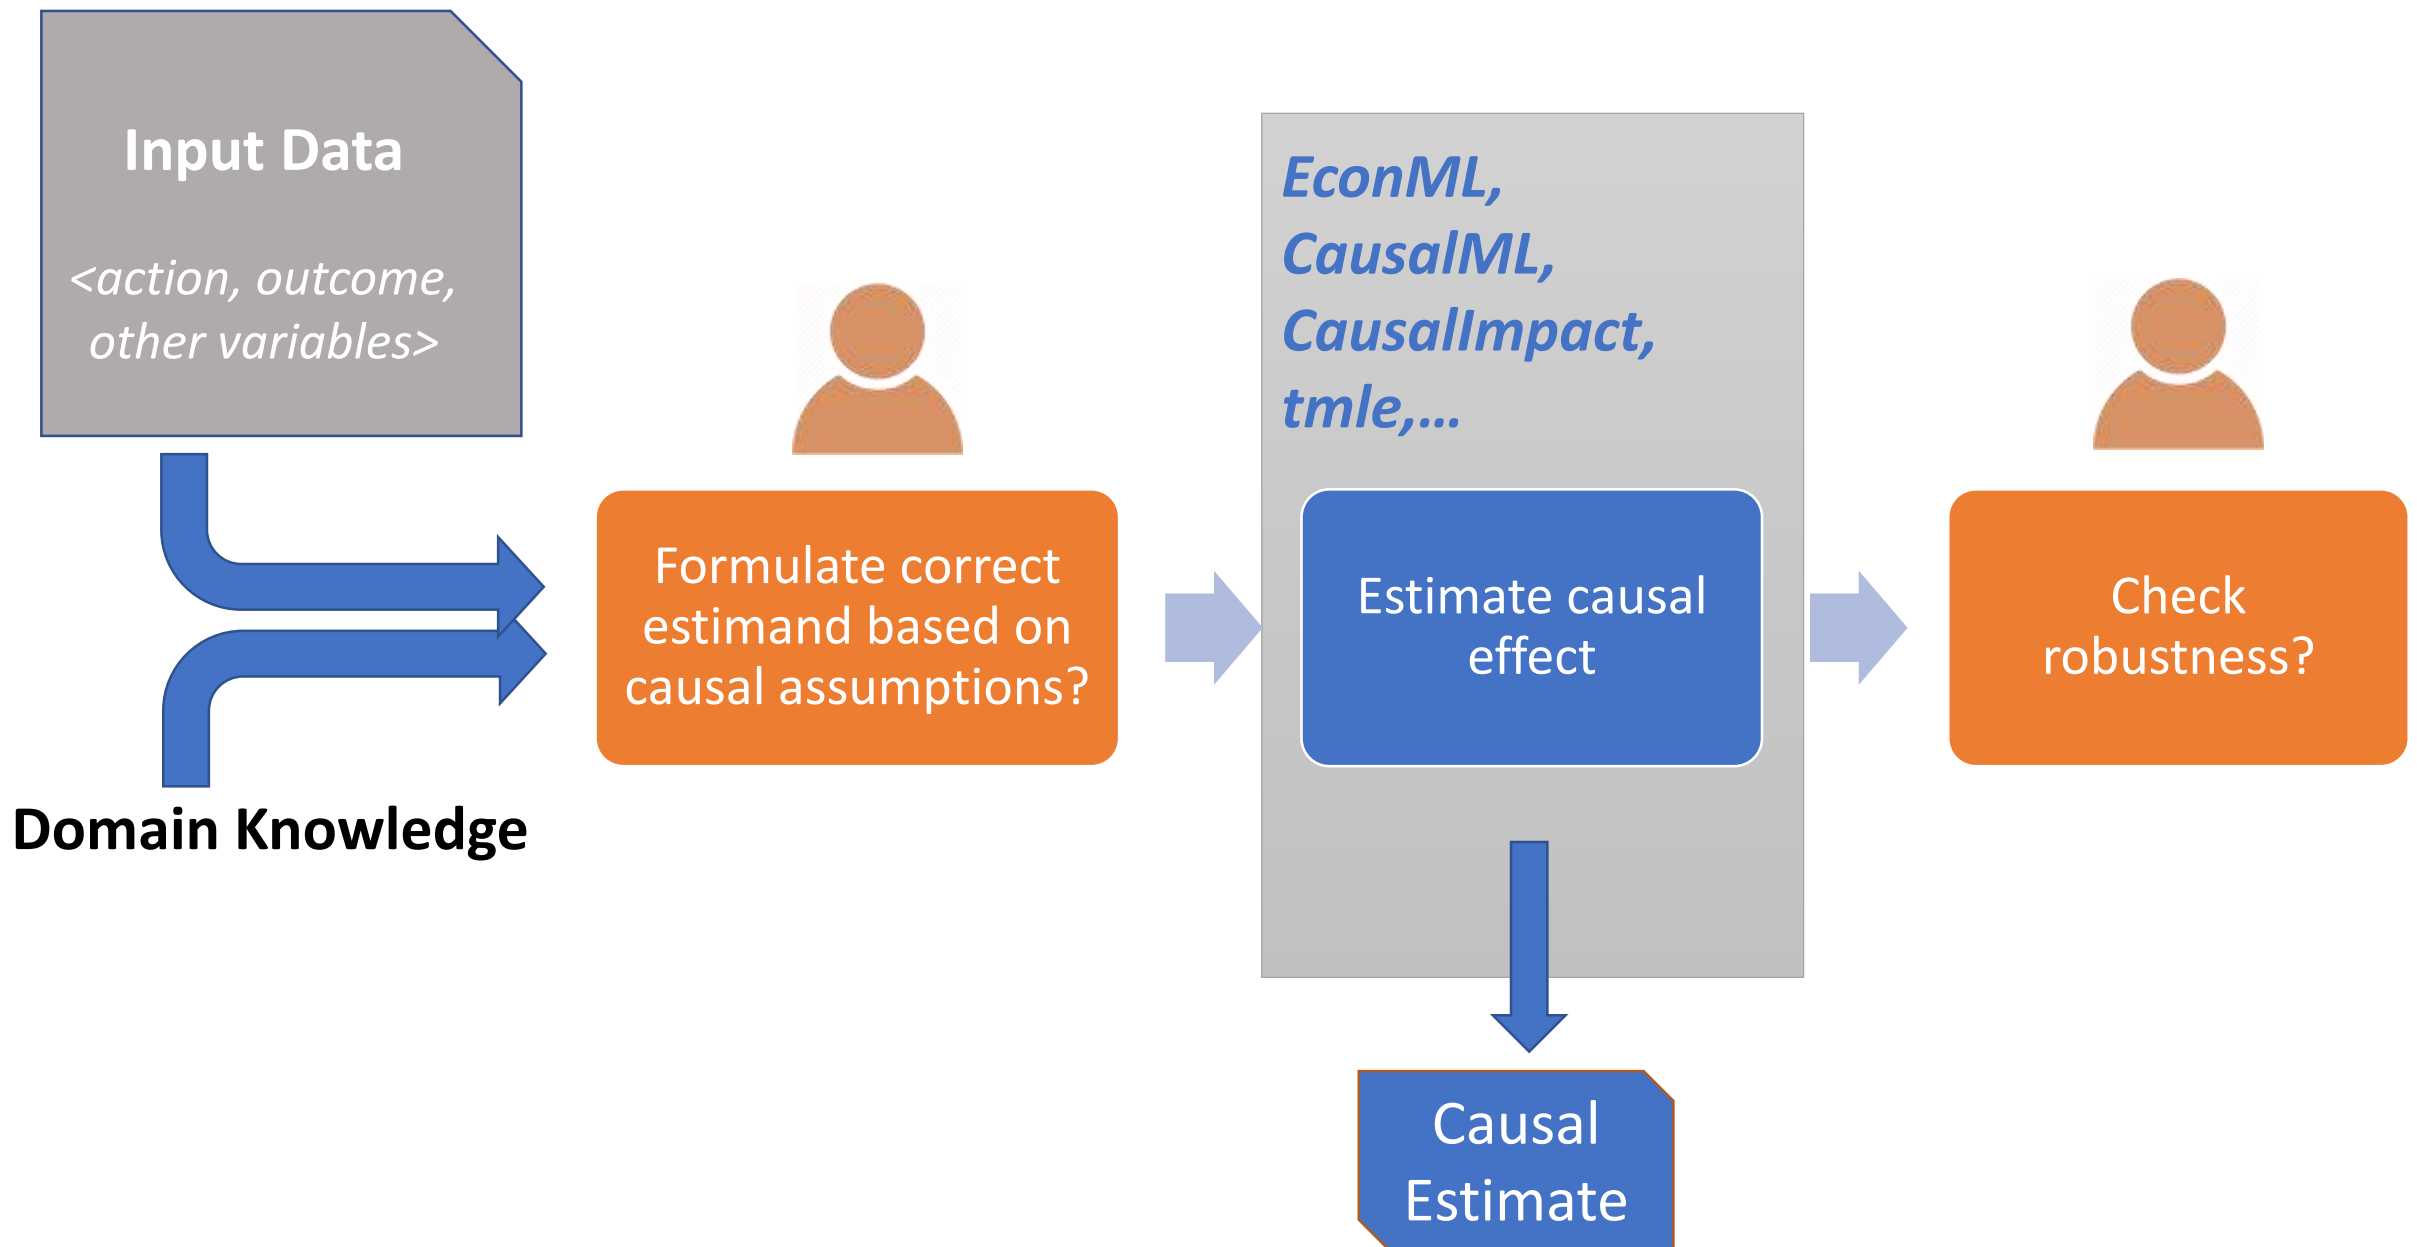
We built [DoWhy library](#) to make assumptions front-and-center of any causal analysis.

- Transparent declaration of assumptions
- Evaluation of those assumptions, to the extent possible

[Most popular](#) causal library on GitHub (>2.4k stars, 300+ forks)

Taught in third-party tutorials and courses: [O'Reilly](#), [PyData](#), [Northeastern](#), ...

An end-to-end platform for doing causal inference




DoWhy

Input Data

*<action, outcome,
other variables>*

Domain Knowledge



Model causal mechanisms

- Construct a causal graph based on domain knowledge

Identify the target estimand

- Formulate correct estimand based on the causal model

Estimate causal effect

- Use a suitable method to estimate effect

Refute estimate

- Check robustness of estimate to assumption violations

Causal effect

DoWhy provides a general API for the four steps of causal inference


1. **Modeling:** Create a causal graph to encode assumptions.
2. **Identification:** Formulate what to estimate.
3. **Estimation:** Compute the estimate.
4. **Refutation:** Validate the assumptions.

We'll discuss the four steps and show a code example using DoWhy.

I. Model the assumptions using a causal graph

Convert domain knowledge to a formal model of **causal assumptions**

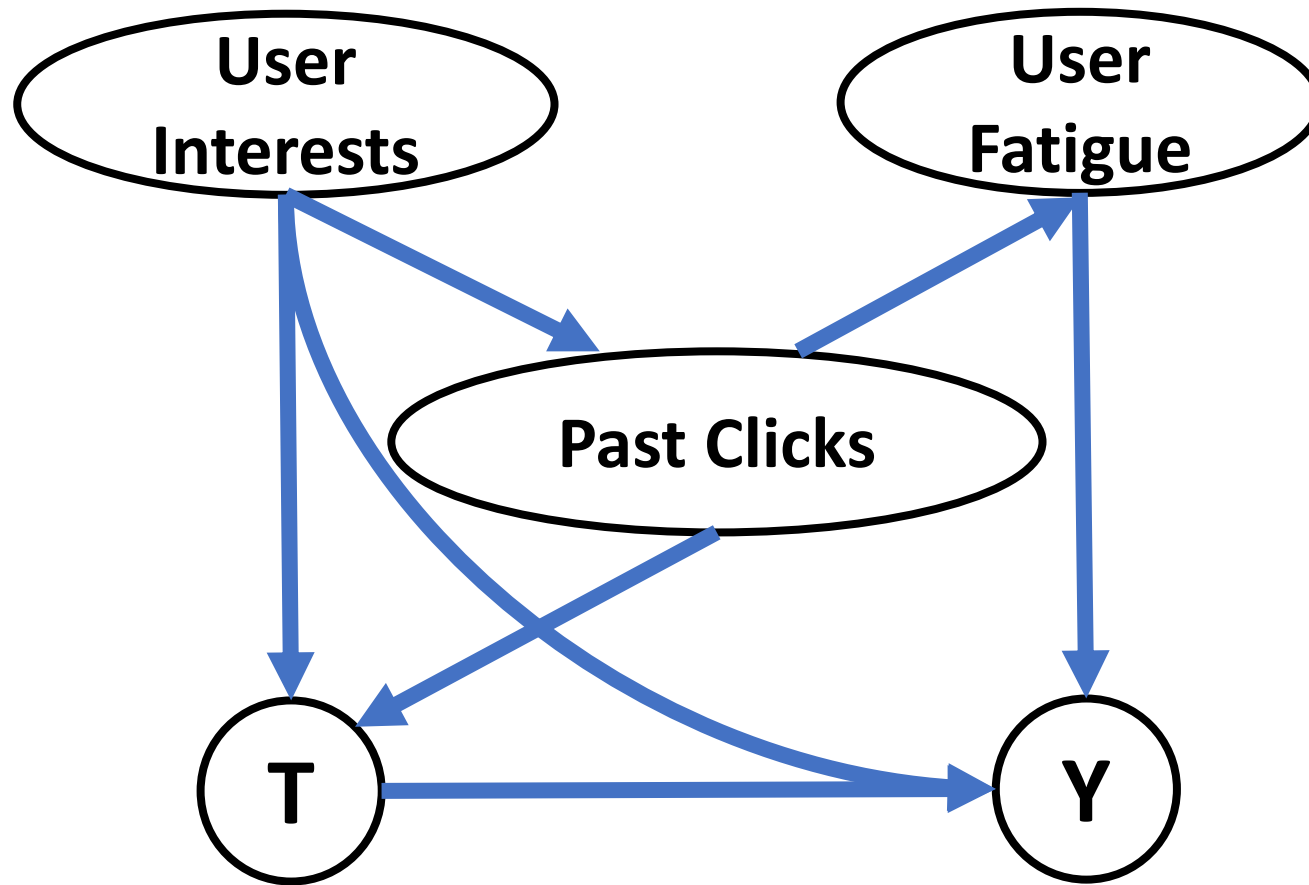
- $A \rightarrow B$ or $B \rightarrow A$?
- Causal graph implies conditional statistical independences
 - E.g., $A \perp\!\!\!\perp C$, $D \perp\!\!\!\perp A \mid B$, ...
 - Identified by *d-separation* rules [Pearl 2009]
- These assumptions significantly impact the causal estimate we'll obtain.



Key intuitions about causal graphs

- Assumptions are encoded by *missing edges*, and *direction* of edges
- Relationships represent stable and independent mechanisms
- Graph cannot be learnt from data alone
- Graphs are a tool to help us reason about a specific problem

Example Graph



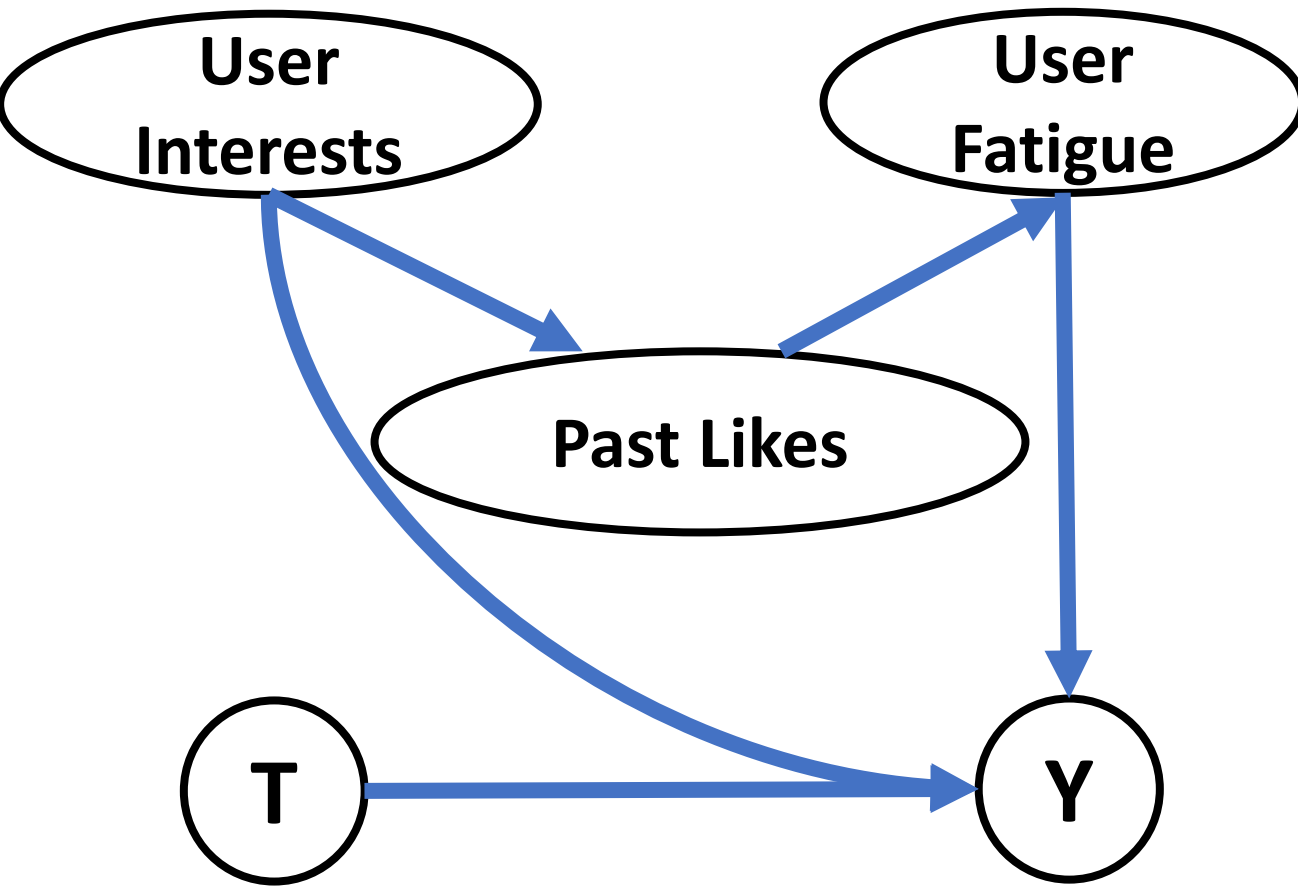
Assumption 1: User fatigue does not affect user interests

Assumption 2: Past clicks do not directly affect outcome

Assumption 3: Treatment does not affect user fatigue.

..and so on.

Intervention is represented by a new graph



Intervention graph:


All edges to Treatment T removed, *keeping everything else the same.*

Represents new data distribution, referred as $do(T)$


Causal effect: $P(Y|do(T))$

II. Identification: Formulate desired quantity and check if it is estimable from given data

**Observed data generated
by this graph**




**Want to answer questions about data that
will be generated by intervention graph**



How to represent quantities from right hand graph (e.g., $P(Y|do(T))$) using only statistical observations from data generated from left hand graph?

Trivial Example: Randomized Experiments

- Observed graph is same as intervention graph in randomized experiment!
 - Treatment T is already generated independent of all other features
 - $\rightarrow P(Y|do(T)) = P(Y|T)$
- **Intuition:** Generalize by simulating randomized experiment
 - When treatment T is caused by other features, Z , adjust for their influence to simulate a randomized experiment



Adjustment Formula and Adjustment Sets

Adjustment formula

$$p(Y|do(T)) = \sum_Z p(Y|T, Z)p(Z)$$

Where Z must be a valid adjustment set:

- The set of all parents of T
- Features identified via *backdoor criterion*
- Features identified via “towards necessity” criterion

Intuitions:

- The union of all features is *not* necessarily a valid adjustment set
- Why not always use parents? Sometimes parent features are unobserved

Many kinds of identification methods

Graphical constraint-based methods

- Randomized and natural experiments
- Adjustment Sets
 - Backdoor, “towards necessity”
- Front-door criterion
- Mediation formula

Identification under additional non-graphical constraints

- Instrumental variables
- Regression discontinuity
- Difference-in-differences

Many of these methods can be used through DoWhy.

III. Estimation: Compute the causal effect

Estimation **uses observed data** to compute the target probability expression from the Identification step.

For common identification strategies using adjustment sets,

$$E[Y|do(T = t), W = w] = E[Y|T = t, W = w]$$


assuming W is a valid adjustment set.

- For binary treatment,

$$\text{Causal Effect} = E[Y|T = 1, W = w] - E[Y|T = 0, W = w]$$

Goal: Estimating conditional probability $Y|T=t$ when all confounders W are kept constant.

Simple Matching: Match data points with the same confounders and then compare their outcomes



Control



Treatment (Cycling)

Simple Matching: Match data points with the same confounders and then compare their outcomes


Identify pairs of treated (j) and untreated individuals (k) who are similar or identical to each other.

$$\text{Match} := \text{Distance}(W_j, W_k) < \epsilon$$

- Paired individuals have almost the same confounders.

Causal Effect =

$$\sum_{(j,k) \in \text{Match}} (y_j - y_k)$$



Challenges of building a good estimator

- **Variance:** If we have a stringent matching criterion, we may obtain very few matches and the estimate will be unreliable.
- **Bias:** If we relax the matching criterion, we obtain many more matches but now the estimate does not capture the target estimand.
- **Uneven treatment assignment:** If very few people have treatment, leads to both high bias and variance.

Need better methods to navigate the **bias-variance tradeoff**.

Machine learning methods can help find a better match for each data point

Synthetic Control: If a good match does not exist for a data point, can we create it synthetically?

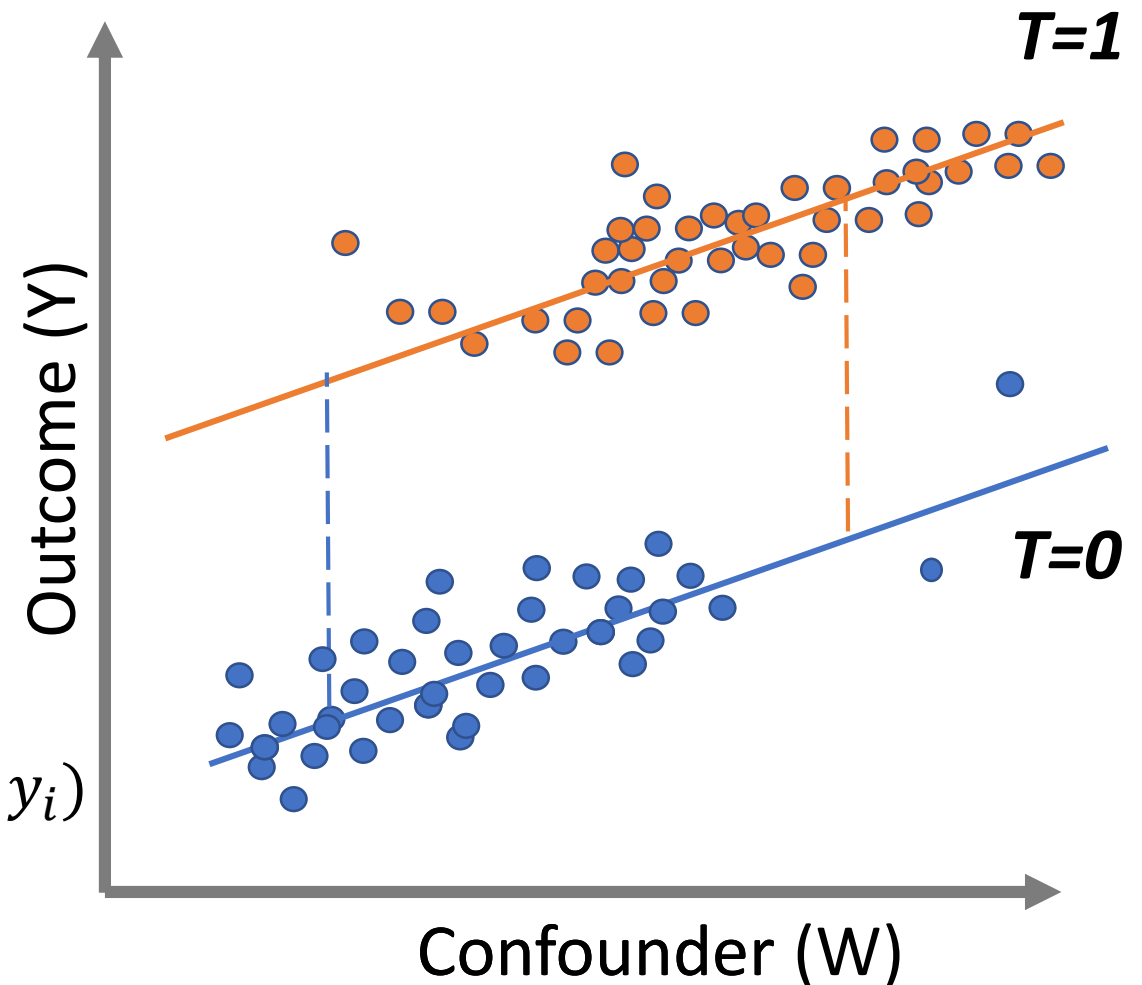
Learn $y = f_{t=0}(w)$,

$y = f_{t=1}(w)$

Assuming f approximates the true relationship between Y and W ,

Causal Effect =

$$\sum_i t_i(y_i - f_{t=0}(w_i)) + (1 - t_i)(f_{t=1}(w_i) - y_i)$$



Other ML methods generalize estimation to continuous treatments

The standard predictor, $y = f(t, w) + \epsilon$ may not provide the right estimate for $\frac{\partial y}{\partial t}$.

Double-ML [Chernozhukov et al. 2016]:

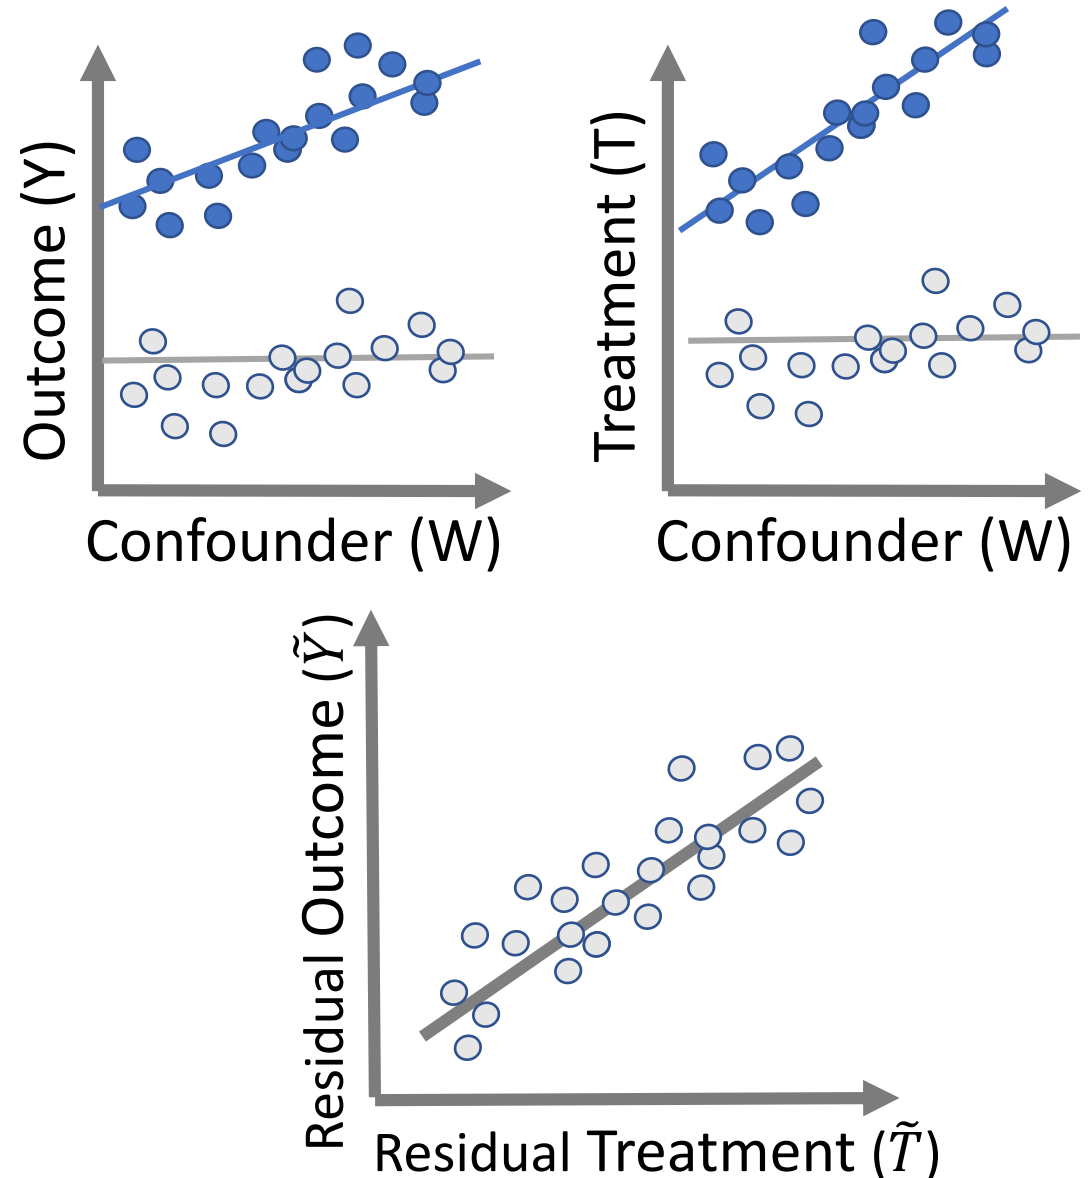
- **Stage 1:** Break down conditional estimation into two prediction sub-tasks.

$$\begin{aligned}\hat{y} &= g(w) + \tilde{y} \\ \hat{t} &= h(w) + \tilde{t}\end{aligned}$$

\tilde{y} and \tilde{t} refer to the unconfounded variation in Y and T respectively after conditioning on w .

- **Stage 2:** A final regression of \tilde{y} on \tilde{t} gives the causal effect.

$$\tilde{y} \sim \beta \tilde{t} + \epsilon$$



Depending on the dataset properties, different estimation methods can be used

Simple Conditioning

- Matching
- Stratification

Propensity Score-Based [Rubin 1983]

- Propensity Matching
- Inverse Propensity Weighting

Synthetic Control [Abadie et al.]

Outcome-based

- Double ML [Chernozhukov et al. 2016]
- T-learner
- X-learner [Kunzel et al. 2017]

Loss-Based

- R-learner [Nie & Wager 2017]

Threshold-based

- Difference-in-differences

All these methods can be called through DoWhy.
(directly or through the Microsoft EconML library)

IV. Robustness Checks: Test robustness of obtained estimate to violation of assumptions

Obtained estimate depends on many (untestable) assumptions.

Model:

Did we miss any unobserved variables in the assumed graph?

Did we miss any edge between two variables in the assumed graph?

Identify:

Did we make any parametric assumption for deriving the estimand?

Estimate:

Is the assumed functional form sufficient for capturing the variation in data?

Do the estimator assumptions lead to high variance?

Best practice: Do refutation/robustness tests for as many assumptions as possible

UNIT TESTS

Model:

- Conditional Independence Test

Identify:

- D-separation Test

Estimate:

- Bootstrap Refuter
- Data Subset Refuter

INTEGRATION TESTS

Test all steps at once.

- Placebo Treatment Refuter
- Dummy Outcome Refuter
- Random Common Cause Refuter
- Sensitivity Analysis
- Simulated Outcome Refuter
/Synth-validation [Schuler et al. 2017]

All these refutation methods are implemented in DoWhy.


Caveat: They can refute a given analysis, *but cannot prove its correctness.*

Example 1: Conditional Independence Refuter

Through its edges, each causal graph implies certain conditional independence constraints on its nodes. [*d-separation, Pearl 2009*]

Model refutation: Check if the observed data satisfies the assumed model's independence constraints.

- Use an appropriate statistical test for independence [*Heinze-Demel et al. 2018*].
- If not, the model is incorrect.



Conditional Independencies:
 $A \perp\!\!\!\perp B$ $A \perp\!\!\!\perp T | W$ $B \perp\!\!\!\perp T | W$

Example 1: Placebo Treatment (“A/A”) Refuter

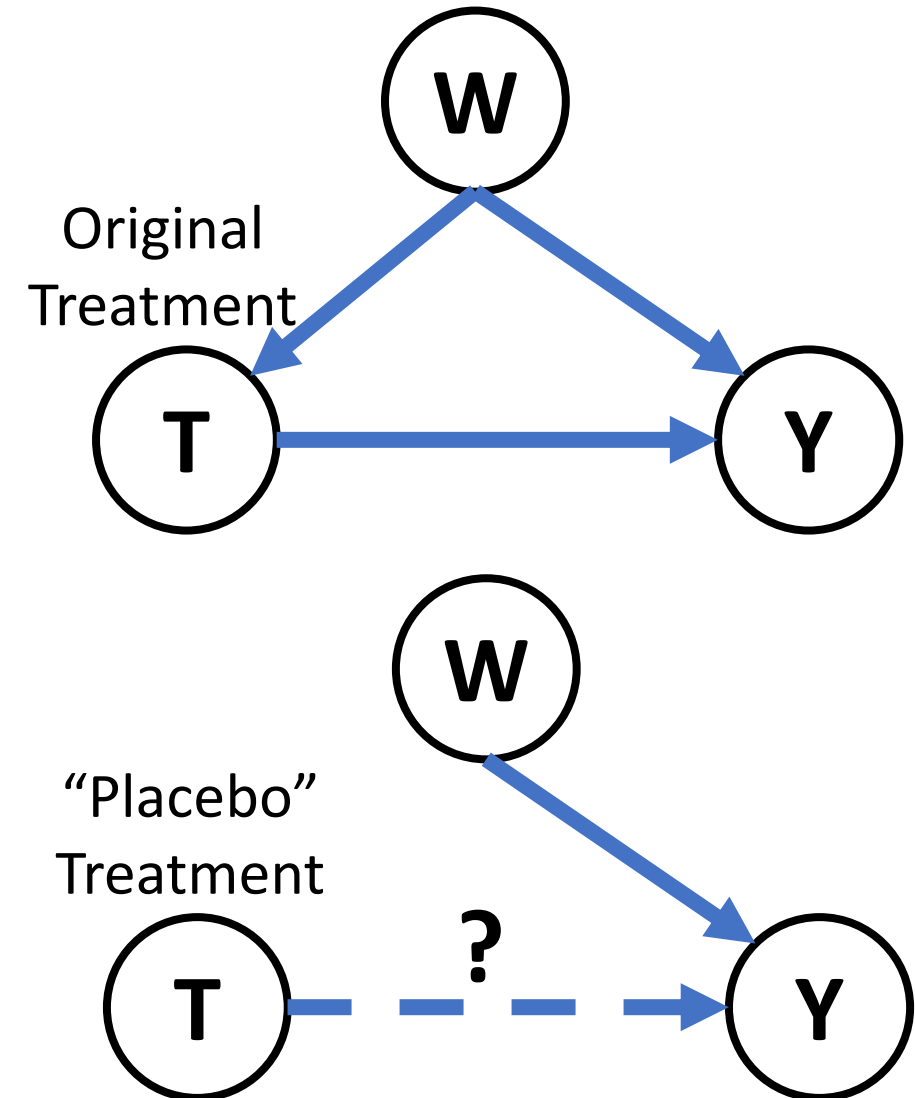
Q: *What if we can generate a dataset where the treatment **does not cause the outcome**?*

Then a correct causal inference method should return an estimate of zero.

Placebo Treatment Refuter:

Replace treatment variable T by a randomly generated variable (e.g., Gaussian).

- Rerun the causal inference analysis.
- If the estimate is significantly away from zero, then analysis is incorrect.




Example 2: Add Unobserved Confounder to check sensitivity of an estimate

Q: *What if there was an unobserved confounder that was not included in the causal model?*

Check how sensitive the obtained estimate is after introducing a new confounder.


Unobserved Confounder Refuter:

- Simulate a confounder based on a given correlation ρ with both treatment and outcome.
 - Maximum Correlation ρ is based on the maximum correlation of any observed confounder.
- Re-run the analysis and check if the sign/direction of estimate flips.



Walk-through of the 4 steps using
the DoWhy Python library

A mystery problem of two correlated variables:
Does x cause y ?



x	y	w
8.329680	16.546904	2.546634
2.083811	4.096492	-3.995819
6.138014	12.041800	0.041479
8.874336	17.833621	2.988497
5.282355	10.481077	-0.860686

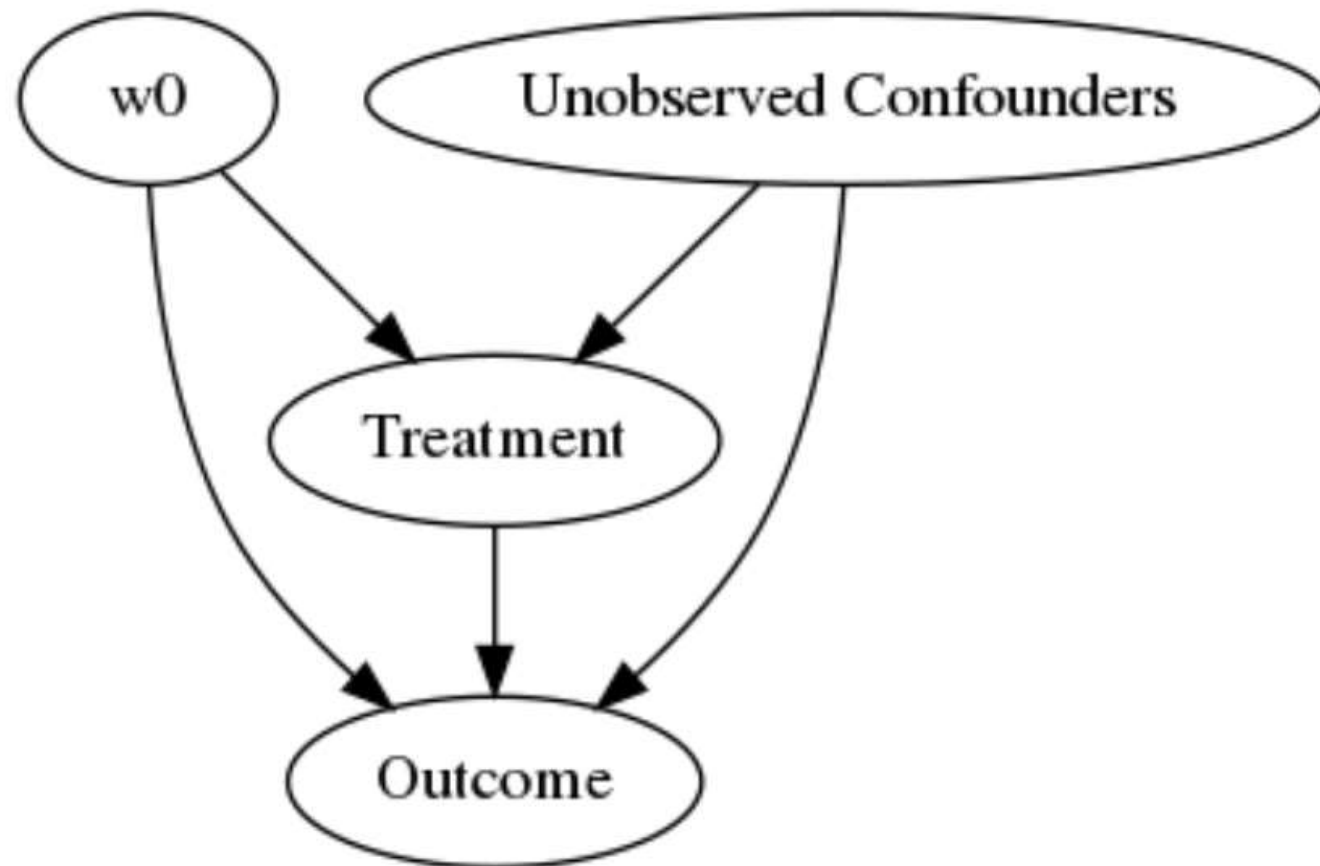
STEP 1: Model the problem as a causal graph

Initializing the causal model.

```
In [5]: model= CausalModel(  
            data=df,  
            treatment=data_dict["treatment_name"],  
            outcome=data_dict["outcome_name"],  
            common_causes=data_dict["common_causes_names"],  
            instruments=data_dict["instrument_names"])  
model.view_model(layout="dot")
```



```
In [6]: from IPython.display import Image, display  
display(Image(filename="causal_model.png"))
```



STEP 2: Identify causal effect using properties of the formal causal graph

Identify the causal effect using properties of the causal graph.

```
In [7]: identified_estimand = model.identify_effect()  
print(identified_estimand)
```

STEP 3: Estimate the causal effect

Once we have identified the estimand, we can use any statistical method to estimate the causal effect.

Let's use Linear Regression for simplicity.

```
In [8]: estimate = model.estimate_effect(identified_estimand,  
      method_name="backdoor.linear_regression")  
print("Causal Estimate is " + str(estimate.value))
```

Step 4: Refuting the estimate

We can also refute the estimate to check its robustness to assumptions (*aka* sensitivity analysis, but on steroids).

Replacing treatment with a random (placebo) variable

```
In [11]: res_placebo=model.refute_estimate(identified_estimand, estimate,  
      method_name="placebo_treatment_refuter", placebo_type="permute")  
print(res_placebo)
```

```
INFO:dowhy.causal_estimator:INFO: Using Linear Regression Estimator
```

```
INFO:dowhy.causal_estimator:b: Outcome~placebo+w0
```

```
Refute: Use a Placebo Treatment
```

```
Estimated effect:(1.0154712956668286,)
```

```
New effect:(-0.001212143363314766,)
```

You can try out this example on Github:

https://github.com/microsoft/dowhy/blob/master/docs/source/example_notebooks/dowhy_confounder_example.ipynb

What else is in DoWhy?

- A unified, extensible API for causal inference that allows external implementations for the 4 steps
 - Can use estimation methods from external libraries such as **EconML** and **CausalML**.

```
dml_estimate = model.estimate_effect(identified_estimand,  
                                     method_name="backdoor.econml.dml.DMLCateEstimator",  
                                     target_units = lambda df: df["X0"]>1,  
                                     confidence_intervals=True,
```

- A convenient CausalDataFrame (contributed by Adam Kelleher)
 - Pandas DataFrame extension with inbuilt functions for calculating causal effect.

Summary: DoWhy, a library that focuses on causal assumptions and their validation

Goal: A unified API for causal inference problems, just like PyTorch or Tensorflow for predictive ML.

Growing open-source community: > 30 contributors

- Roadmap: More powerful refutation tests, counterfactual prediction.
- Please contribute! Would love to hear your ideas on Github.

Resources

- DoWhy Library: <https://github.com/microsoft/dowhy>
- Arxiv paper on the four steps: <https://arxiv.org/abs/2011.04216>
- Upcoming book on causality and ML: <http://causalinference.gitlab.io/>

thank you– Amit Sharma
(@amt_shrma)